


(a) Cross Platform Inference Latency



(b) A100 Throughput vs Batch Size

